

# Contributions to Deep Transfer Learning

## From Supervised to Reinforcement Learning

**Matthia Sabatelli**

Montefiore Institute, Department of Electrical Engineering and  
Computer Science, Université de Liège, Belgium

March 30th 2022

# Presentation outline

## ① Part I: Preliminaries

Deep Transfer Learning

## ② Part II: Supervised Learning

On the Transferability of Convolutional Neural Networks

On the Transferability of Lottery Winners

## ③ Part III: Reinforcement Learning

A Novel Family of Deep Reinforcement Learning Algorithms

On the Transferability of Deep-Q Networks



# PART I

# Deep Transfer Learning

## PART II

We start by defining the components of **Supervised Learning**:

- An input space  $\mathcal{X}$ ,
- An output space  $\mathcal{Y}$ ,
- A joint probability distribution  $P(X, Y)$ .
- A loss function  $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$

The **goal** is to find a function  $f : \mathcal{X} \rightarrow \mathcal{Y}$  that minimizes the expected risk:

$$R(f) = \mathbb{E}_{(\vec{x}, y) \sim P(X, Y)} [\ell(y, f(\vec{x}))]$$

# On the Transferability of Convolutional Neural Networks

Typically, the only information available to build this function  $f$  is a learning sample of input-output pairs

$LS_T = \{(x_i, y_i) | i = 1, \dots, N_T\}$  drawn independently from  $P_T(x, y)$ .

However, when it comes to [Transfer Learning](#) we have an additional dataset  $LS_S$  that can be used for finding a better  $f$  than when only  $LS_T$  is used for training.

# On the Transferability of Convolutional Neural Networks

For our first study ...

- We consider  $f$  to come in the form of a **Convolutional Neural Network**
- We define the source domain  $\mathcal{D}_S$  to be that of **natural images**
- We define the target domain  $\mathcal{D}_T$  to be that of **Digital Heritage**
- We assume that labels are available in both the source and target data and that the input spaces  $\mathcal{X}_T$  and  $\mathcal{X}_S$  match:  
**Inductive Transfer Learning**

# On the Transferability of Convolutional Neural Networks

We have a CNN pre-trained on  
the ImageNet dataset as  
source-task  $\mathcal{T}_S$



We have three target-tasks  $\mathcal{T}_T$   
defined by the Rijksmuseum  
collection

# On the Transferability of Convolutional Neural Networks



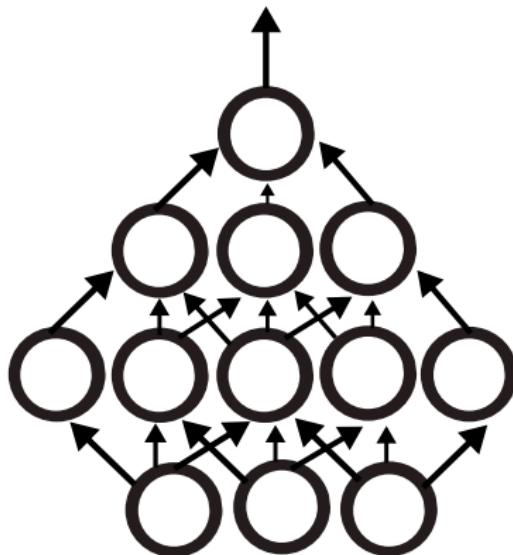
# On the Transferability of Lottery Winners

**The Lottery Ticket Hypothesis:** "*A randomly-initialized dense neural network contains a subnetwork that is initialized such that when trained in isolation it can match the test accuracy of the original network after training for at most the same number of iterations.*<sup>1</sup>"

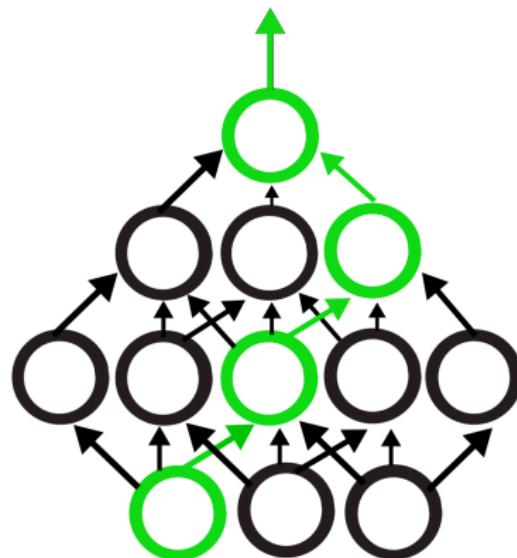
---

<sup>1</sup>Frankle, Jonathan, and Michael Carbin. "The lottery ticket hypothesis: Finding sparse, trainable neural networks." (2018)

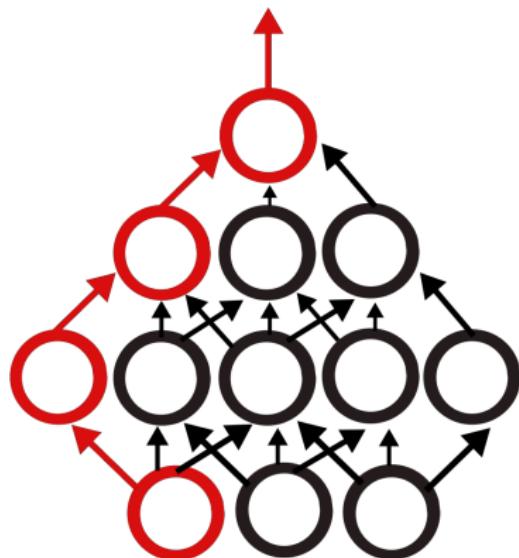
## On the Transferability of Lottery Winners



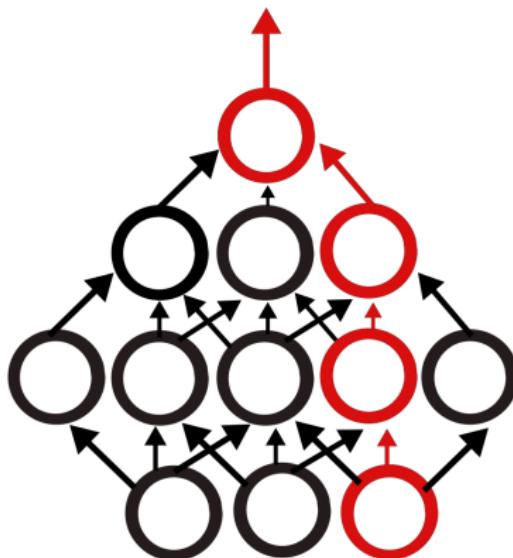
# On the Transferability of Lottery Winners



# On the Transferability of Lottery Winners



# On the Transferability of Lottery Winners



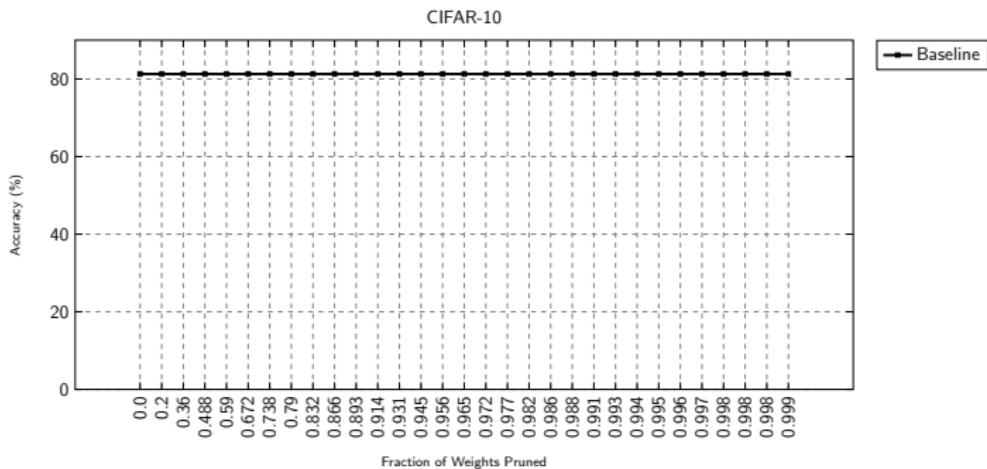
# On the Transferability of Lottery Winners

⇒ How does one **find** a winning ticket?

- Randomly initialize a network  $f(x; \theta_0)$  where  $\theta_0 \sim \mathcal{D}_\theta$
- We train the network for  $j$  iterations
- We prune  $p\%$  of the parameters in  $\theta_j$  creating a mask  $m$
- Reset the remaining parameters to their values at  $\theta_0$ ,  
creating a winning ticket  $f(x; m \odot \theta_0)$

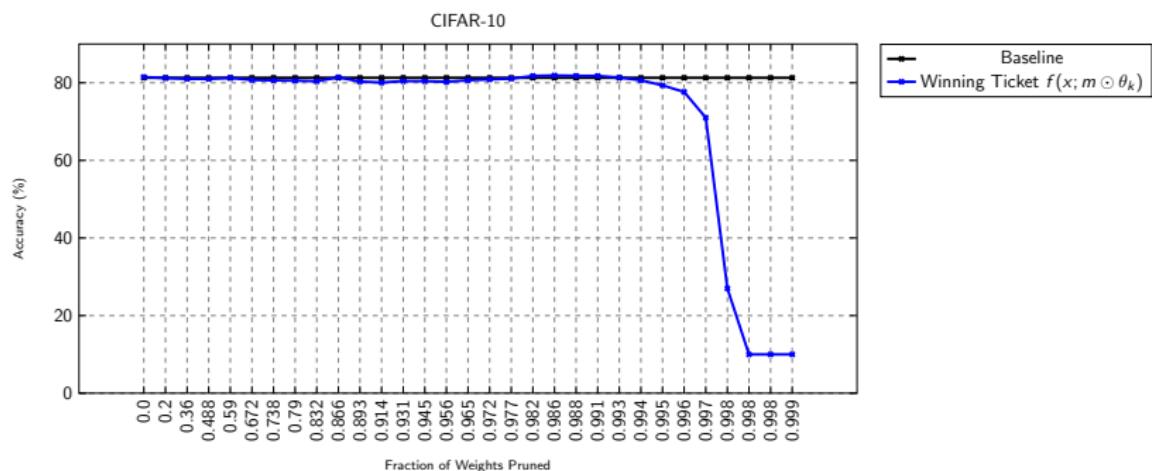
# On the Transferability of Lottery Winners

## The Lottery Ticket Hypothesis in practice ...



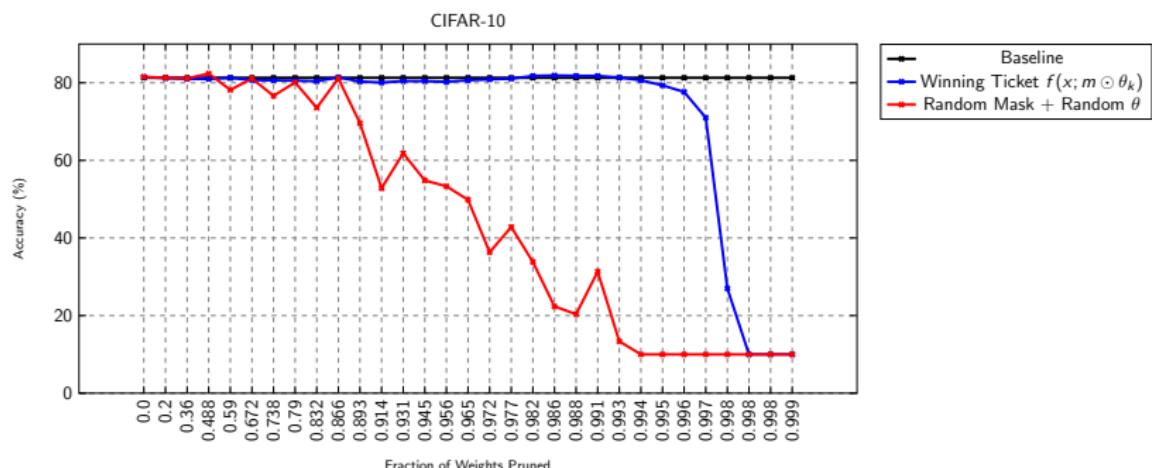
# On the Transferability of Lottery Winners

The Lottery Ticket Hypothesis in practice ...



# On the Transferability of Lottery Winners

The Lottery Ticket Hypothesis in practice ...



# On the Transferability of Lottery Winners

Why are lottery tickets  $f(x; m \odot \theta_0)$  so **special**?

- Train faster
- Faster Inference
- (Sometimes) obtain a better final performance

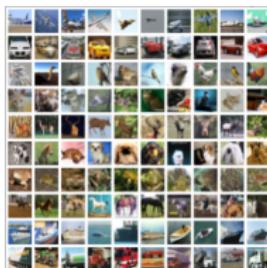
However ...

- Identifying a winning ticket is **computationally expensive**
- We **do not** know how and why lottery winners appear throughout learning

# On the Transferability of Lottery Winners

⇒ Therefore we study whether winning tickets  $f(x; m \odot \theta_0)$  found on natural image datasets can get transferred to the non-natural realm

We use three popular Computer Vision datasets as source domains  $\mathcal{D}_S$ :



# On the Transferability of Lottery Winners

⇒ Therefore we study whether winning tickets  $f(x; m \odot \theta_0)$  found on natural image datasets can get transferred to the non-natural realm

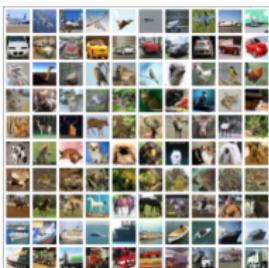
We use three popular Computer Vision datasets as source domains  $\mathcal{D}_S$ :



# On the Transferability of Lottery Winners

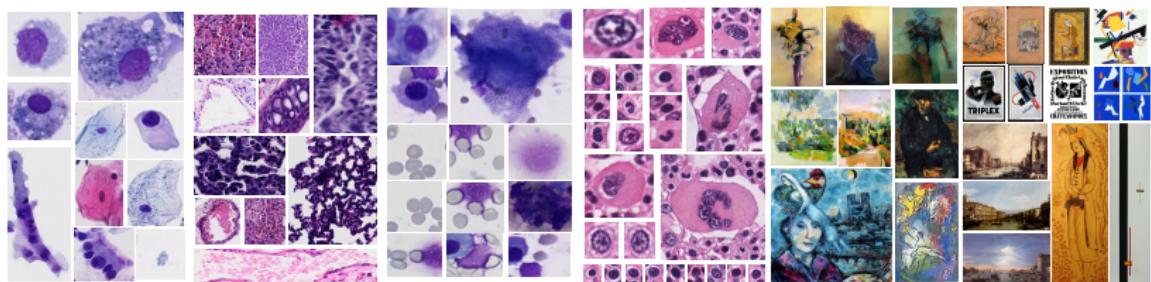
⇒ Therefore we study whether winning tickets  $f(x; m \odot \theta_0)$  found on natural image datasets can get transferred to the non-natural realm

We use three popular Computer Vision datasets as source tasks  $\mathcal{T}_S$



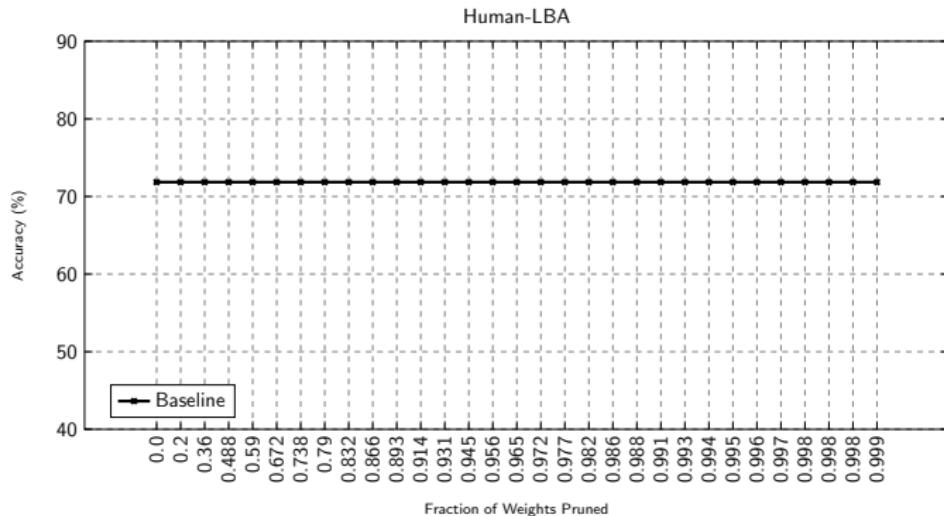
# On the Transferability of Lottery Winners

And we use **seven datasets** of non-natural images as target tasks  $\mathcal{T}_T$  coming from the fields of Digital Pathology and Digital Heritage



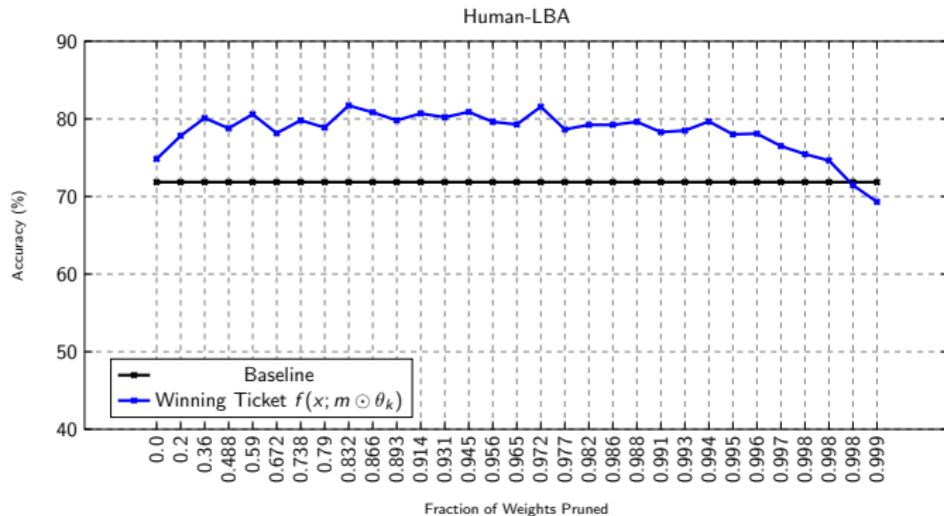
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



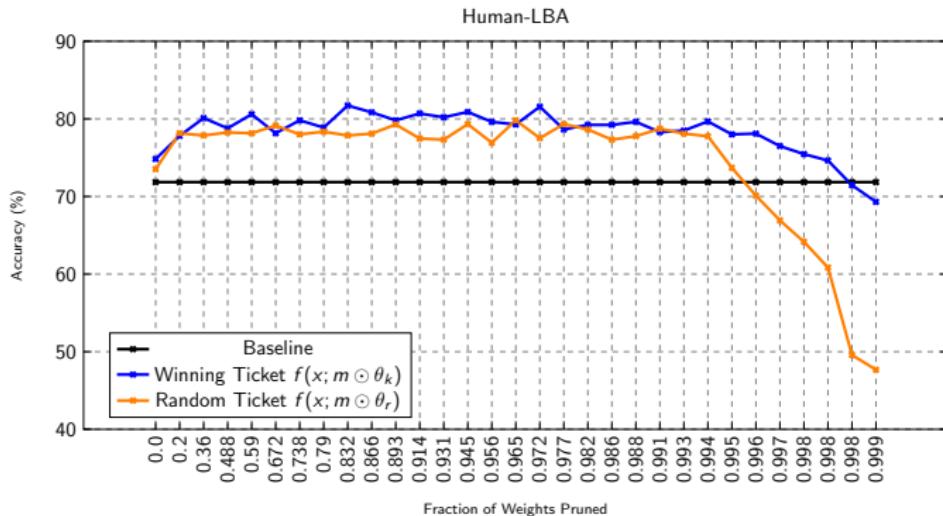
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



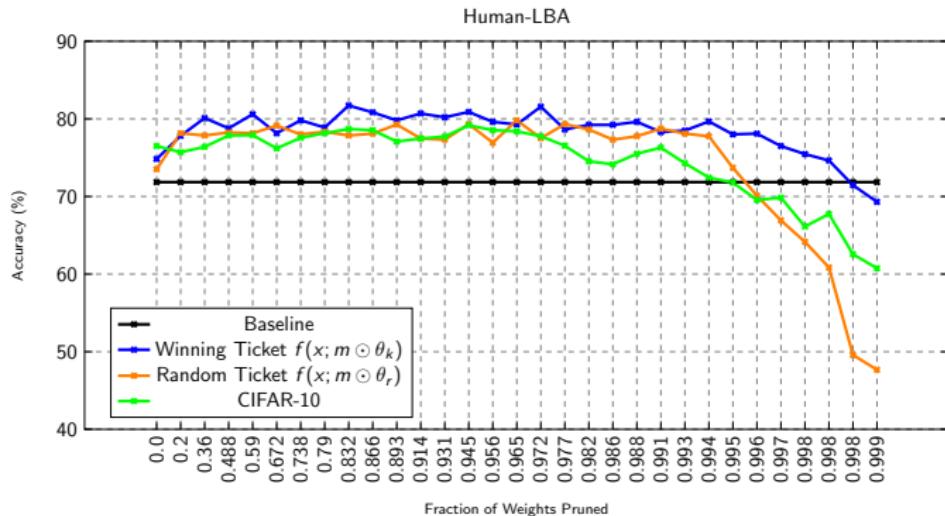
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



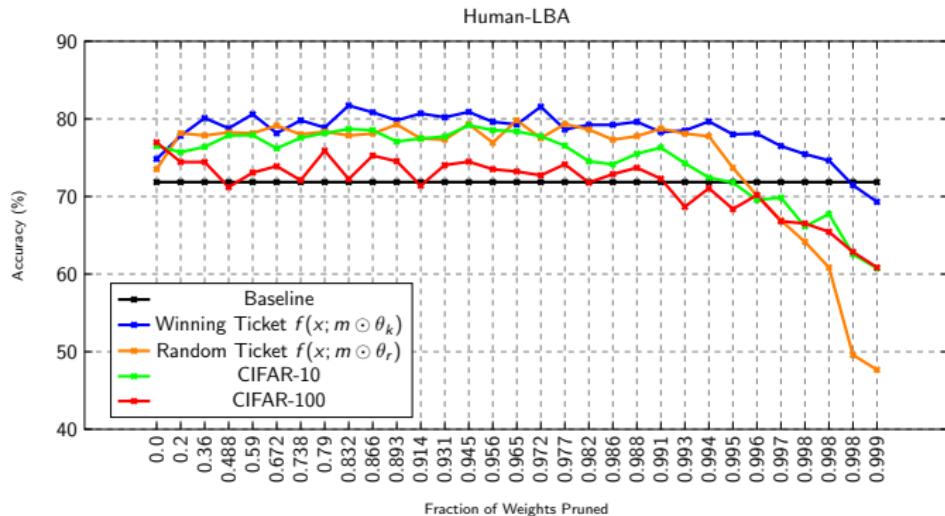
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



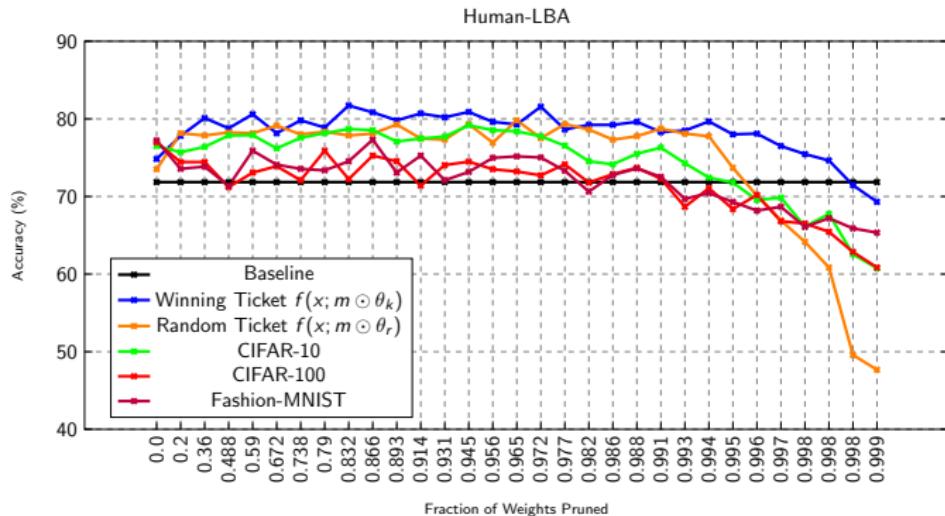
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



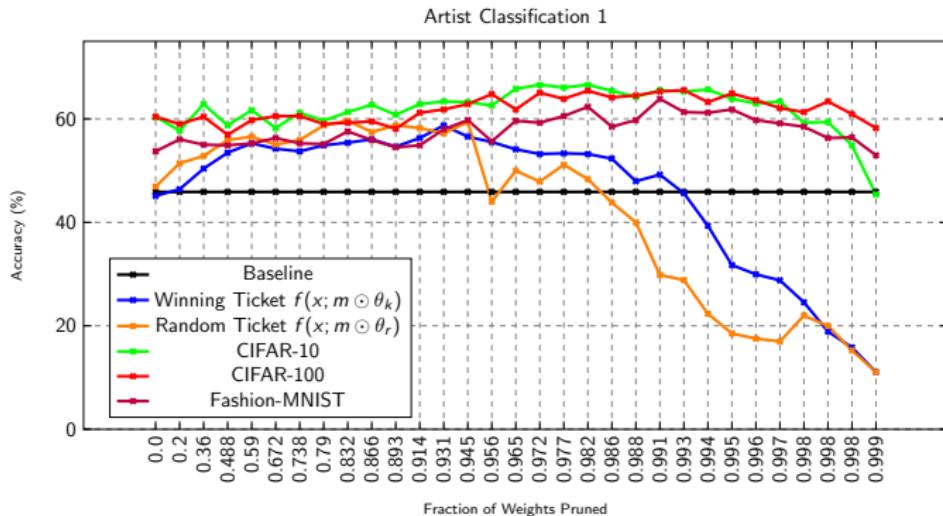
# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Pathology](#) data ...



# On the Transferability of Lottery Winners

Transferring lottery winners to the [Digital Heritage](#) data ...



# On the Transferability of Lottery Winners

Our main findings show that on Digital Pathology data

- All lottery winners significantly outperform unpruned models
- Natural lottery winners contain a generic inductive bias (to some extent)
- Best performance is obtained by identifying a winning ticket directly on the target task  $\mathcal{T}_T$

whereas on Digital Heritage data

- Natural lottery winners transfer much better
- They can even outperform target task  $\mathcal{T}_T$  tickets

## On the Transferability of Lottery Winners

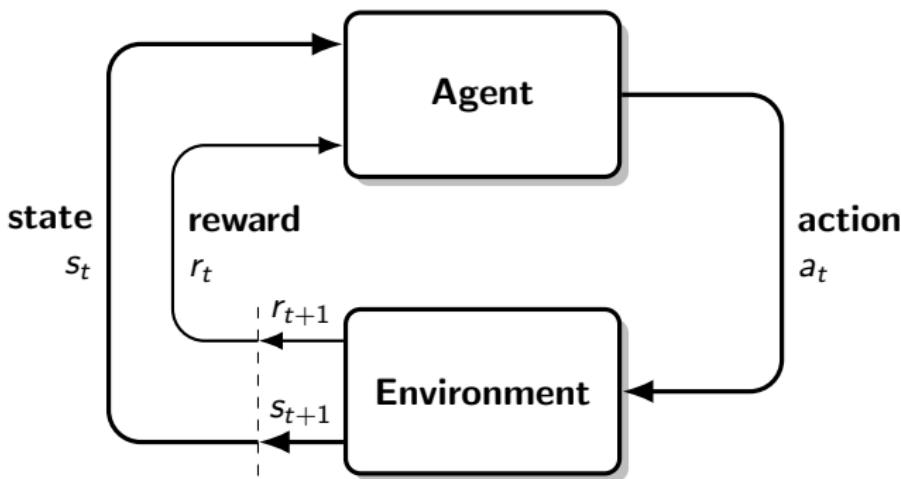
We also provide additional empirical insights into the Lottery Ticket Hypothesis:

- We show that completely pre-trained winning tickets  $f(x; m \odot \theta_i)$  **overfit** on the source task  $\mathcal{T}_S$
- The presence of winning tickets **does not depend** on the size of the training data
- The closer the source task  $\mathcal{S}_T$  and the target task  $\mathcal{T}_T$  the **better** the transferability of  $f(x; m \odot \theta_k)$

## PART III

# A Novel Family of Deep Reinforcement Learning Algorithms

We now consider a different machine learning paradigm:  
**Reinforcement Learning** (RL), where an agent needs to learn how to interact with its environment



# A Novel Family of Deep Reinforcement Learning Algorithms

Such interaction is modeled as a **Markov Decision Process** (MDP) consisting of the following elements:

- A set of possible states  $\mathcal{S}$ ,
- A set of possible actions  $\mathcal{A}$ ,
- A transition function  $\mathcal{P} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow [0, 1]$ ,
- A reward function  $\mathfrak{R} : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \rightarrow \mathbb{R}$ ,
- A discount factor denoted as  $\gamma \in [0, 1]$ .

Therefore we can define an MDP as  $\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathfrak{R}, \gamma \rangle$ .

# A Novel Family of Deep Reinforcement Learning Algorithms

The interaction between the agent and the environment is given by the agent's **policy**  $\pi$ , a probability distribution over  $a \in \mathcal{A}(s)$  for each  $s \in \mathcal{S}$ :

$$\pi(a|s) = \Pr \{a_t = a | s_t = s\}, \text{ for all } s \in \mathcal{S} \text{ and } a \in \mathcal{A}.$$

The **goal** of the agent is to maximize the expected discounted return as:

$$\begin{aligned} G_t &= r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots \\ &= \sum_{k=0}^{\infty} \gamma^k r_{t+k}. \end{aligned}$$

# A Novel Family of Deep Reinforcement Learning Algorithms

In RL some components of the MDP are unknown

$$\mathcal{M} = \langle \mathcal{S}, \mathcal{A}, \mathcal{P}, \mathcal{R}, \gamma \rangle,$$

which in the case of **model-free** RL can be overcome by learning either the state-value function

$$V^\pi(s) = \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k} \middle| s_t = s, \pi \right],$$

or the state-action value function

$$Q^\pi(s, a) = \mathbb{E} \left[ \sum_{k=0}^{\infty} \gamma^k r_{t+k} \middle| s_t = s, a_t = a, \pi \right].$$

# A Novel Family of Deep Reinforcement Learning Algorithms

- Model-free algorithms are implemented in a **tabular** fashion, meaning that the state values, or state-action values, are stored within tables of sizes  $|\mathcal{S}|$  and  $|\mathcal{S} \times \mathcal{A}|$  respectively
- For many problems we typically seek to learn an **approximation** of the value functions

$$V^\pi(s) \approx V^\pi(s; \theta) \text{ and } Q^\pi(s, a; \theta) \approx Q(s, a; \theta)$$

# A Novel Family of Deep Reinforcement Learning Algorithms

This approximation can be represented by a Convolutional Neural Network

# A Novel Family of Deep Reinforcement Learning Algorithms

Typical DRL algorithms aim at only learning an approximation of the state-action value function  $Q^\pi(s, a) \approx Q^\pi(s, a; \theta)$

- A policy  $\pi$  can only be derived from  $Q^\pi(s, a)$
- It is hard to combine RL algorithms with neural networks
- Dealing with one value function is enough

However ...

- Training can be very slow
- Algorithms are prone to diverge
- Do not correctly estimate  $Q^\pi(s, a)$

# A Novel Family of Deep Reinforcement Learning Algorithms

⇒ Therefore we suggest to **jointly approximate** the state-value function  $V^\pi(s; \phi)$  alongside the state-action value function  $Q^\pi(s, a; \theta)$

We can do this by either learning  $V^\pi(s; \phi)$  with

$$L(\phi) = \mathbb{E}_{\langle s_t, a_t, r_t, s_{t+1} \rangle \sim U(D)} \left[ (r_t + \gamma V(s_{t+1}; \Phi^-) - V(s_t; \Phi))^2 \right],$$

and  $Q^\pi(s, a)$  with:

$$L(\theta) = \mathbb{E}_{\langle s_t, a_t, r_t, s_{t+1} \rangle \sim U(D)} \left[ (r_t + \gamma V(s_{t+1}; \Phi^-) - Q(s_t, a_t; \theta))^2 \right],$$

# A Novel Family of Deep Reinforcement Learning Algorithms

... or by learning  $V^\pi(s; \phi)$  with

$$L(\phi) = \mathbb{E}_{(s_t, a_t, r_t, s_{t+1}) \sim U(D)} \left[ (r_t + \gamma \max_{a \in \mathcal{A}} Q(s_{t+1}, a; \theta^-) - V(s_t; \Phi))^2 \right],$$

and  $Q^\pi(s, a)$  with:

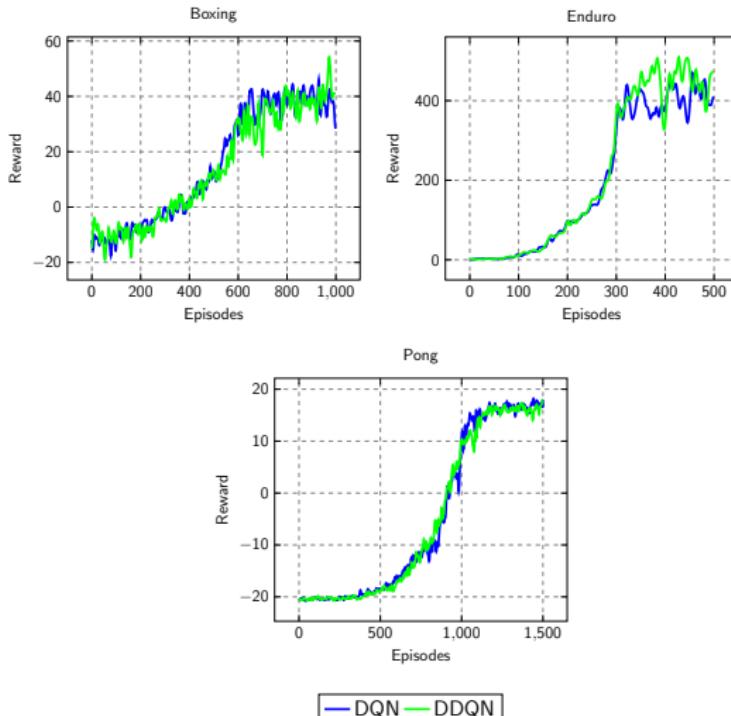
$$L(\theta) = \mathbb{E}_{(s_t, a_t, r_t, s_{t+1}) \sim U(D)} \left[ (r_t + \gamma V(s_{t+1}; \Phi) - Q(s_t, a_t; \theta))^2 \right].$$

# A Novel Family of Deep Reinforcement Learning Algorithms

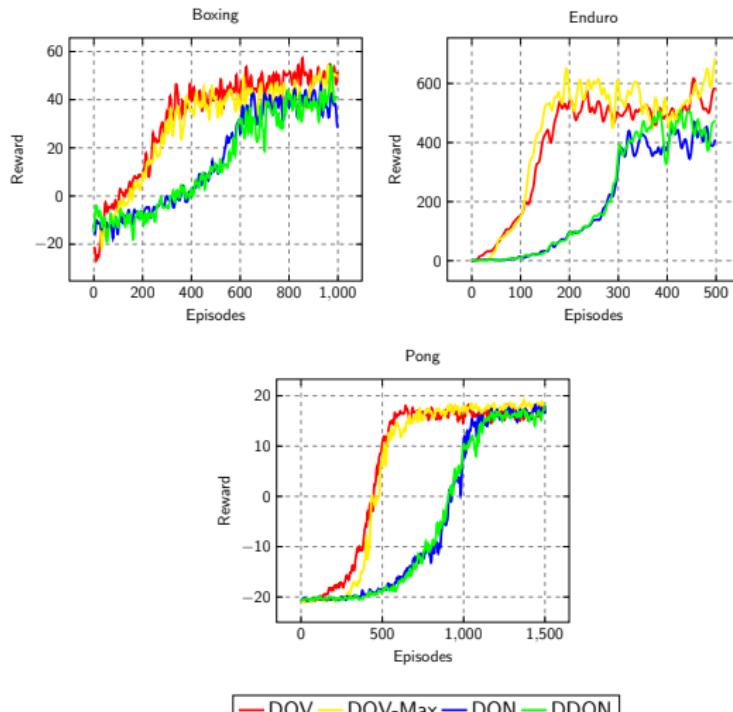
The first two objective functions constitute the **DQV-Learning** algorithm, whereas the second two objective functions define the **DQV-Max Learning** algorithm.

We compare their performance to algorithms which only learn an approximation of the state-action value function  $Q^\pi(s, a)$

# A Novel Family of Deep Reinforcement Learning Algorithms

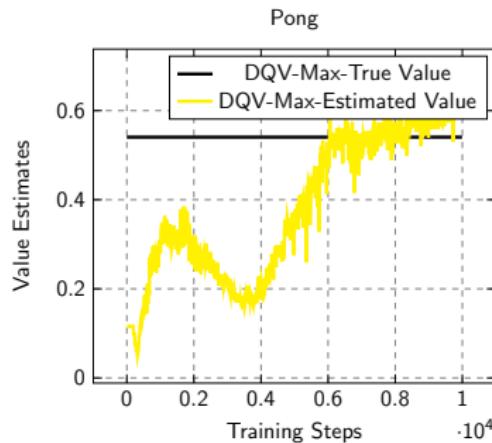
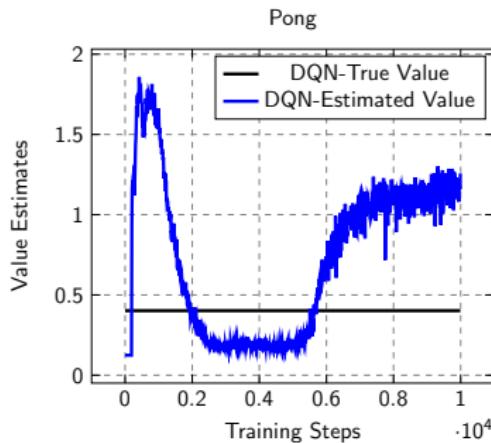


# A Novel Family of Deep Reinforcement Learning Algorithms



# A Novel Family of Deep Reinforcement Learning Algorithms

⇒ Jointly approximating two value functions over one results in faster convergence, but also in accurate value estimates:



# A Novel Family of Deep Reinforcement Learning Algorithms

To conclude DQV-Learning and DQV-Max Learning

- Result in **faster learning**
- Are **robust** to the "Deadly Triad of DRL"
- Suffer **less** from the overestimation bias of the  $Q$  function
- **Scale well** to the multi-agent setting <sup>2</sup>

However ...

- Memory wise are **twice** as expensive as DQN and DDQN
- **Do not** always result in better policies

---

<sup>2</sup>Leroy, Pascal, et al. "QVMix and QVMix-Max: Extending the Deep Quality-Value Family of Algorithms to Cooperative Multi-Agent Reinforcement Learning." (2021).



## Final References